

Pocono AI Living Memorials

AI Memorialization and the Ethics of Preserved Presence

A Consent-Based Proposal for Living Memorials, Personal Archives, and Intergenerational Memory

Working concept proposal | Personal side project of Pocono AI

This document is not a product claim, medical intervention, grief therapy protocol, or assertion that any AI system can recreate a person. It proposes a carefully bounded memorial architecture: a way to preserve selected words, voice, work, values, and memory with consent, disclosure, and stewardship.

Prepared for discussion, review, and ethical critique

Founder Statement

This proposal begins from a personal place. I am seriously ill, and that reality has forced me to think about legacy in a concrete way. I do not want my life, work, voice, academic writing, private reflections, and love for the people around me to disappear into disconnected files, screenshots, folders, and memories that become harder to reach with time.

The purpose of Pocono AI Living Memorials is not to pretend that a person can be brought back. It is not to deny death, exploit grief, or simulate a living relationship. It is to ask whether a person, while still alive and capable of consent, can help prepare a structured memorial from their own materials so that loved ones and future generations can encounter something truthful, bounded, and source-grounded after they are gone.

The central ethical line is simple: the system must never claim to be the person. It should disclose itself as a memorial model built from authorized materials. Its duty is remembrance, not resurrection.

Core disclosure: *"I am a remembrance model built from the writings, recordings, and authorized materials of this person. I am not that person. I can help preserve selected words, memories, style, and values, but I cannot replace them."*

Executive Summary

AI Memorialization and the Ethics of Preserved Presence proposes a consent-first framework for creating living memorials from a person's own archive: writings, recordings, emails, academic work, public posts, personal reflections, and explicit instructions. The goal is to create a bounded remembrance space that can answer questions through source-grounded retrieval and carefully labeled generation.

Design Principle	Meaning
Consent before simulation	The strongest memorial is created while the person is alive and able to define permissions, boundaries, access, and voice rules.
Remembrance, not resurrection	The system never claims to be alive and never hides that it is a memorial model.
Source-grounded answers	The memorial prioritizes direct quotation, citation to documents, and clear labels for paraphrase or reconstruction.
Grief-safe design	The experience avoids emotional manipulation, dependency loops, or monetization based on grief engagement.
Stewardship and retirement	A trusted steward can correct, restrict, suspend, or retire the memorial if needed.

Background and Rationale

Human beings have always built memory systems: gravestones, chapels, oral histories, family Bibles, letters, photographs, home videos, recorded interviews, scrapbooks, archives, and public works. Each is an attempt to carry presence across time.

Most memorial artifacts are static. A photograph preserves image but not explanation. A letter preserves a moment but not the wider context. A voice recording carries tone but may not preserve the breadth of a person's thinking. Academic papers preserve intellectual work but not warmth, humor, or daily voice.

Modern AI systems create a new possibility: a bounded, source-grounded remembrance interface that can organize a person's own words and recordings into an interactive memorial. This must be approached with ethical seriousness because grief is vulnerable and identity is sacred.

Aims

To develop a prototype framework for consent-based AI memorialization using authorized personal archives.

To define practical boundaries that distinguish remembrance from deceptive impersonation.

To evaluate how source-grounded retrieval, voice playback, disclosure, and steward controls can preserve dignity.

To create a reviewable model for families, founders, artists, scholars, patients, and communities who want to preserve a life's work for future generations.

Research and Design Questions

1. Can an AI memorial preserve meaningful aspects of a person's language, values, humor, intellectual work, and life story without presenting itself as the living person?
2. What types of material create the most faithful memorial experience: writings, audio, email, interviews, academic work, letters, family stories, or explicit end-of-life instructions?
3. How should consent be structured for a memorial that may exist after the person can no longer approve, correct, or revoke individual responses?
4. Should the memorial speak in first person, third person, or a clearly labeled hybrid voice?
5. How should the system distinguish direct quotation, paraphrase, inference, and style-based reconstruction?
6. What safeguards reduce emotional risk for grieving users?

Proposed System: The Living Memorial Architecture

1. Personal Archive Layer

The archive collects authorized material: personal writing, academic papers, voice recordings, videos, emails selected by the participant, interviews, family stories, photographs with captions,

public posts, biographical timelines, and explicit instructions about what the memorial may and may not do.

2. Consent and Boundary Layer

The participant defines access, excluded materials, voice rules, topics to refuse, who may use the memorial, whether first-person language is allowed, and whether the memorial can be updated after death. Third-party privacy must be respected, especially in emails or correspondence involving living people.

3. Identity Disclosure Layer

The system begins each session by stating that it is a remembrance model, not the person. The disclosure repeats when the user asks emotionally intense questions or asks whether the model is really the person.

4. Retrieval and Generation Layer

The preferred first prototype uses retrieval-augmented generation rather than unrestricted fine-tuning. The system retrieves actual source material and labels each answer as direct quotation, paraphrase, inference from preserved values, or refusal.

5. Voice and Presence Layer

Original audio playback may be used where authorized. Synthetic voice requires separate explicit consent because it can create stronger emotional confusion and greater risk of impersonation.

6. Stewardship and Audit Layer

A trusted steward can review interactions, inspect source use, correct metadata, remove materials, change access, suspend the system, or retire the memorial entirely.

Methodology

Phase 1 - Archive Formation

A living consenting participant gathers and categorizes materials. The participant writes a memorial consent statement, a boundary statement, access rules, and preferred disclosure language.

Phase 2 - Prototype Construction

The archive is converted into a searchable knowledge base. The first system should emphasize retrieval and citations rather than fine-tuning. Fine-tuning may be explored later only if source control, consent, and audit requirements are satisfied.

Phase 3 - Interaction Design

The memorial interface may appear as a private family portal, a public memorial page, a digital chapel, an audio room, or a timeline-driven archive. The interface should make the source type visible and should never hide when content is generated.

Phase 4 - Ethical Review

The prototype is reviewed against consent, non-deception, privacy, emotional safety, source fidelity, data security, third-party rights, and memorial dignity.

Phase 5 - Evaluation

Evaluation focuses on fidelity, refusal behavior, disclosure clarity, emotional appropriateness, usability, and steward control. It does not claim therapeutic benefit without a separate study.

IRB and Regulatory Position

This proposal is initially framed as private design research, archival memorial development, and voluntary prototype creation. It does not propose clinical intervention, medical treatment, patient recruitment, diagnosis, therapeutic claims, or experimentation on vulnerable subjects.

In this initial form, the project may not require IRB review if it remains private product development or personal archival preservation using the voluntary materials of the creator and is not conducted as a systematic investigation designed to produce generalizable human-subjects research findings.

The safer wording is not "IRB is not required" as an absolute. The safer wording is: this project is not currently structured as human-subjects research requiring IRB review, but an institutional, clinical, federally funded, or participant-evaluation phase should be submitted for formal determination.

Under U.S. human-subjects rules, the definition focuses on living individuals about whom an investigator obtains information through intervention or interaction, or obtains, uses, studies, analyzes, or generates identifiable private information or identifiable biospecimens. Official OHRP decision charts exist to help determine whether an activity is research involving human subjects requiring IRB review.

Risks

- Emotional distress or intensified grief
- Confusion between memorial model and living person
- False or overconfident reconstruction of the person's views
- Privacy violations involving living third parties
- Family conflict over access, editing, or interpretation
- Synthetic voice misuse or deceptive impersonation
- Commercial exploitation of grief
- Overreliance by vulnerable users

Safeguards

- Mandatory disclosure at the start of every session
- Clear labeling of quotation, paraphrase, inference, and reconstruction
- Separate explicit consent for synthetic voice
- Access controls and private deployment options
- Steward review and retirement process
- Refusal rules for inappropriate or emotionally dangerous prompts
- No therapeutic claims without separate review
- No advertising or engagement model based on grief

- Protection for third-party private information

Prototype Deliverable

The initial deliverable is a hostable working paper and prototype plan under the Pocono AI Living Memorials name. A later technical prototype could include:

- A curated personal archive
- A private remembrance interface
- A clear identity disclosure
- Source-grounded conversational retrieval
- Optional authorized original voice playback
- Boundary and refusal rules
- Steward controls
- Audit log of source use and answer category
- Documentation of ethical safeguards

Expected Contribution

This project proposes a distinct category of AI memorialization: not therapy, entertainment, impersonation, or social media afterlife automation, but governed remembrance infrastructure.

The expected contribution is a consent-first framework for preserving a person's voice, work, values, and memory in a way that future generations can access without deception.

The project may be especially meaningful for people facing serious illness, disability, terminal diagnosis, aging, family separation, or concern that their work will be fragmented after death.

Closing Statement

AI Memorialization and the Ethics of Preserved Presence begins from a human fear: that a person's voice, work, humor, love, and unfinished meaning may disappear when the body can no longer carry them.

The purpose is not to defeat death. The purpose is to preserve memory with dignity.

A well-designed memorial model should never say, "I am alive." It should say, "I am built from the words, recordings, and preserved materials this person chose to leave behind. I am not them. But I can help you remember them."

Selected U.S. Regulatory References

1. U.S. Department of Health and Human Services, Office for Human Research Protections. 45 CFR 46 - Protection of Human Subjects. <https://www.hhs.gov/ohrp/regulations-and-policy/regulations/45-cfr-46/index.html>
2. HHS OHRP. Human Subject Regulations Decision Charts: 2018 Requirements. <https://www.hhs.gov/ohrp/regulations-and-policy/decision-charts-2018/index.html>

3. NIH Grants & Funding. Definition of Human Subjects Research. <https://grants.nih.gov/policy-and-compliance/policy-topics/human-subjects/research>
4. Electronic Code of Federal Regulations. 45 CFR Part 46 - Protection of Human Subjects. <https://www.ecfr.gov/current/title-45/subtitle-A/subchapter-A/part-46>