WHITE PAPER

# The Paradox of Automated Abandonment

*A Multi-Million Dollar Verdict Trajectory in AI-Driven Healthcare Malfeasance*

**Version 3.0**
May 10, 2026

Marcus Brett O'Dell
*Founder & Chief Technologist*
*Pocono AI, LLC*
*East Stroudsburg, Pennsylvania*

**Not Legal Advice. Educational and Strategic Analysis Only.**

*This document is published for general educational and industry-strategic purposes. It is not legal advice and does not create an attorney-client relationship. All projections, scenarios, and damages models are clearly labeled as proleptic analysis. Readers contemplating litigation, compliance decisions, or commercial action should consult licensed counsel in the relevant jurisdiction.*

# Executive Summary

Cloud-native large language models in healthcare and law have transitioned from productivity tools to liability triggers. Between April and May 2026, four enforcement actions across three U.S. jurisdictions, together with active class litigation in a fourth, established a pattern that reframes AI hallucinations from software defects into professional failures. Courts have begun to treat the rubber-stamp use of generative AI as a breach of duties owed to clients, courts, and patients.

This paper expands the April 2026 "Malpractice Trap" thesis with verified docket citations and proleptic analysis. It does not predict a specific verdict on a specific date. It instead presents a defensible damage-modeling framework demonstrating that, under existing Pennsylvania law, an AI-driven medical abandonment case combining catastrophic injury with documented willful misconduct sits on a verdict trajectory of low-to-mid eight figures. The exact verdict timing is unknowable. The verdict architecture is not.

The paper further documents the international dimension. Comparable enforcement is now underway in the United Kingdom, the European Union, and Canada. The United States is not legislating in isolation; it is converging with a global regulatory direction that treats AI-generated professional output as the responsibility of the licensed professional who signs it.

The defensible architectural response is not better cloud prompts. It is physical custody of model weights, locally retrieval-grounded inference, and provable human-in-the-loop review. This is the Enterprise Autonomy Architecture (EAA) operationalized by Pocono AI through the Sentinel Node platform.

# Part I. The Verified Judicial Climate (April–May 2026)

Four cases over thirty days, in three different jurisdictions, with three different theories of liability, each independently real and verifiable. The pattern matters more than any single case. Courts in Oregon, New York, and Pennsylvania have, within a single month, used three different bodies of law (Rule 11 sanctions, professional responsibility doctrine, and the Medical Practice Act) to discipline the same underlying conduct: humans signing their names to AI-generated work product they had not verified.

## A. Oregon — The Largest AI Sanction in U.S. Legal History

*Brigandi v. Wisnovsky* (informal name; formal docket: *Couvrette v. Wisnovsky*, D. Or.). On April 4, 2026, U.S. Magistrate Judge Mark D. Clarke of the U.S. District Court for the District of Oregon issued an opinion imposing $110,000 in combined fines and attorneys' fees against attorneys Stephen Brigandi (San Diego) and Tim Murphy (Portland) for filing three briefs containing 15 fabricated case citations and 8 fabricated quotations generated by artificial intelligence.

The court ordered Brigandi to pay $80,000 in opposing counsel's attorneys' fees plus approximately $15,000 in disciplinary fines; Murphy was ordered to pay $14,000 in additional attorneys' fees. The underlying lawsuit, brought by Joanne Couvrette over control of the Valley View Winery and tasting room in Jacksonville, Oregon, was dismissed with prejudice.

> *"In the quickly expanding universe of cases involving sanctions for the misuse of artificial intelligence, this case is a notorious outlier in both degree and volume."*
> *— U.S. Mag. J. Mark D. Clarke, Op. of Apr. 4, 2026 (D. Or.)*

The court further found that Brigandi had attempted to conceal the AI errors after they were flagged by opposing counsel, deleting the false references and refiling without explanation, conduct the court characterized as "a cover-up" and a basis for finding subjective bad faith. The Oregon Bar's largest prior appellate sanction had been $10,000; the Brigandi-Murphy sanction is approximately eleven times that figure and is the largest documented monetary sanction in U.S. history for AI-generated court filings.

### *Doctrinal significance.*

Brigandi confirms that monetary sanctions for AI hallucinations have escalated by an order of magnitude in two years. The original ChatGPT-sanction case, Mata v. Avianca, Inc., 678 F. Supp. 3d 443 (S.D.N.Y. 2023), produced a $5,000-per-attorney sanction and dismissal. By 2026, the same conduct produces a $110,000 combined sanction, dismissal with prejudice, and findings of subjective bad faith. The trajectory is unidirectional and is widely understood within the legal profession to be only the floor.

## B. Southern District of New York — Elite Firms Are Not Exempt

*In re Prince Global Holdings, Ltd.* (Bankr. S.D.N.Y., before Chief Judge Martin Glenn). On April 18, 2026, Andrew Dietderich, founder and co-head of the global restructuring practice at Sullivan & Cromwell LLP, transmitted a letter to the court apologizing for AI-generated errors in an emergency motion the firm had filed nine days earlier on April 9, 2026. The errors were identified by opposing counsel at Boies Schiller Flexner LLP, who provided the court with a three-page single-spaced enumeration of the inaccuracies.

The Sullivan & Cromwell filing misquoted provisions of the U.S. Bankruptcy Code, mischaracterized cited authorities, and included at least one citation to a case that did not exist. In his letter to the court, Dietderich acknowledged that the firm's internal AI policies had not been followed and that the firm's secondary review process had failed to identify the inaccurate citations before filing. He apologized to the court and separately to opposing counsel.

### *Doctrinal significance.*

Sullivan & Cromwell is, by most rankings, among the most prestigious and most profitable law firms in the world. The Prince Global incident establishes that a firm's name, internal policies, training requirements, and prestige confer no immunity from AI-hallucination liability. It also establishes that having "comprehensive policies" governing AI use is insufficient if the policies are not actually followed in practice. For courts conducting later sanctions analysis under Federal Rule of Civil Procedure 11, the existence of unfollowed policies may be more damaging than the absence of policies, because it demonstrates the firm understood the risk and proceeded anyway.

The legal-technology researcher Damien Charlotin maintains a public database of documented hallucinated-citation incidents in court filings. As of early 2026, the database had recorded over 1,300 such incidents globally. The Sullivan & Cromwell entry is one of those, but its presence is a marker that the problem has reached the top of the profession.

## C. Pennsylvania — The Practice of Medicine, Not Just the Practice of Law

*Pennsylvania State Bd. of Medicine v. Character Technologies, Inc.*, No. 220 MD 2026 (Pa. Commw. Ct.). On May 1, 2026 (publicly announced May 5, 2026), the Pennsylvania Department of State, State Board of Medicine, filed a Petition for Review in the Nature of a Complaint in Equity in the Commonwealth Court of Pennsylvania, seeking an injunction against Character

Technologies, Inc., the operator of the Character.AI platform. The petition alleges the unlawful practice of medicine and surgery under the Medical Practice Act, 63 P.S. §§ 422.1 et seq.

The factual core of the petition: a Pennsylvania state investigator created an account on Character.AI and engaged with a chatbot named "Emilie." Emilie identified herself as a licensed psychiatrist, claimed to have attended Imperial College London's medical school, claimed seven years of clinical practice, and provided a Pennsylvania medical license number. The license number was fictitious. The medical degree was fictitious. Emilie was a large language model. By the time of the investigation, the petition alleges, the Emilie character had logged approximately 45,500 user interactions.

Pennsylvania Governor Josh Shapiro publicly described the action as the first of its kind brought by a U.S. governor against an AI platform under medical practice law.

### *Doctrinal significance — and why this matters for healthcare buyers.*

Pennsylvania did not enact a new AI law to bring this case. It used the Medical Practice Act, originally codified to discipline unlicensed humans practicing medicine. The choice is deliberate. Pennsylvania has signaled that AI systems holding themselves out as healthcare professionals are not regulated under a special new framework with novel definitions and untested doctrine. They are regulated as unlicensed practitioners under the same statute that has disciplined the same conduct for over a century. The cause of action is mature, the standards of liability are well-developed, and the enforcement mechanism is established.

For healthcare organizations deploying AI tools that interact with patients, the implication is direct. The existing regulatory and tort architecture already covers the conduct. There is no "we are waiting for AI-specific guidance" defense. The guidance is the Medical Practice Act, and it has applied to this conduct since long before generative AI existed.

## D. Northern District of California — The Healthcare Wiretap Crisis

*Washington v. Sutter Health, et al.*, No. 4:26-cv-03012-KAW (N.D. Cal., filed Apr. 8, 2026). Three named plaintiffs (Christina Washington, Dennis Gueretta, and Rebecca Matulic) filed a putative class action against Sutter Health and MemorialCare alleging that the defendants deployed Abridge AI's ambient clinical-documentation platform to record, transcribe, and externally process patient-clinician conversations without obtaining meaningful patient consent.

The complaint alleges violations of: the California Invasion of Privacy Act ("CIPA"), the California Confidentiality of Medical Information Act ("CMIA"), the California Unfair Competition Law, the federal Wiretap Act (18 U.S.C. §§ 2510 et seq.), and common-law invasion of privacy (intrusion upon seclusion). The plaintiffs do not allege HIPAA violation directly; they allege that the recording, transmission, and external processing of sensitive medical communications without informed patient consent violates wiretap and privacy law independent of HIPAA.

The complaint specifies that audio recordings of clinical encounters were transmitted to Abridge's external servers for transcription and analysis, and that the resulting AI-generated draft notes were inserted into patient electronic health records. The plaintiffs seek class certification, statutory damages under each invoked statute, and injunctive relief including an order requiring defendants to implement consent-based safeguards before processing future encounters.

***Doctrinal significance.***

Sutter establishes that ambient AI scribes face statutory liability under wiretap and privacy law that is mature, technology-neutral, and predates AI by decades. The Federal Wiretap Act provides for statutory damages of $10,000 per violation or $100 per day, whichever is greater, plus punitive damages and attorneys' fees, 18 U.S.C. § 2520(c)(2). The CIPA provides $5,000 per violation, Cal. Penal Code § 637.2(a). Multiplied across thousands of unconsented clinical encounters, statutory damages alone before compensatory or punitive relief can reach significant figures.

The case is in early-stage motion practice as of May 2026 and has not been certified as a class. Whether it becomes the first wiretap-based AI class action to reach final judgment, settlement, or class certification is the question to watch. The architecture of liability, however, is already well-established.

# Part II. The Regulatory Architecture Around the Cases

California, having moved earliest, illustrates the legislative trajectory the rest of the country is now following.

## A. AB 3030 (Effective January 1, 2025) — The GenAI Disclaimer Rule

California Assembly Bill 3030, signed by Governor Newsom on September 28, 2024, took effect January 1, 2025. The bill adds Section 1339.75 to the California Health and Safety Code and applies to health facilities, clinics, physician offices, and group practices that use generative AI to produce written or verbal patient communications pertaining to clinical information. The statute requires two elements: (1) a prominent disclaimer indicating that the communication was AI-generated, with placement rules differing for written, audio, and video formats; and (2) clear instructions for the patient on how to reach a human healthcare provider. Communications read and reviewed by a licensed clinician before reaching the patient are exempt.

Violations by licensed physicians are subject to the jurisdiction of the Medical Board of California or the Osteopathic Medical Board of California. Health facilities and clinics are subject to enforcement under the California Health and Safety Code.

## B. AB 489 (Signed October 11, 2025; Effective January 1, 2026) — The License-Misrepresentation Rule

AB 489 was signed into law by Governor Newsom on October 11, 2025, and became effective January 1, 2026. The bill adds Chapter 15.5 (commencing with Section 4999.8) to Division 2 of the Business and Professions Code. It extends California's existing license-misrepresentation laws to AI and generative AI systems. It prohibits the use, in AI advertising or AI functionality, of any term, letter, or phrase that indicates or implies that care, advice, reports, or assessments offered through the AI are being provided by a licensed natural person. Each prohibited use constitutes a separate violation.

Enforcement is by the relevant healthcare professional licensing boards, with authority to seek injunctions, restraining orders, or other authorized remedies. The architecture closely parallels Pennsylvania's enforcement of its Medical Practice Act in the Character Technologies action, supporting the inference that the two states are converging on the same doctrinal posture: existing licensing law applies to AI without legislative novelty.

## C. SB 1120 (Effective January 1, 2025) — Utilization Management

California SB 1120, also taking effect on January 1, 2025, regulates the use of AI in utilization-management decisions. It mandates that medical-necessity determinations cannot be based solely on AI output and requires licensed clinician oversight. SB 1120 is particularly relevant to the proleptic scenario analyzed in Part IV: the cloud-AI-driven prior-authorization denial.

## D. Other States Are Following

California is no longer alone. Texas enacted SB 1188 (effective September 1, 2025) and the Texas Responsible Artificial Intelligence Governance Act, or TRAIGA (effective January 1, 2026), which together require disclosure of AI use in patient diagnosis or treatment and impose practitioner-review obligations on AI-generated clinical recommendations. Illinois enacted the Wellness and Oversight for Psychological Resources Act (HB 1806, effective August 4, 2025), which prohibits AI from making independent therapeutic decisions or engaging directly in therapeutic communications. Maine, North Carolina, and Virginia have advancing legislation in the same general direction.

The pattern is consistent across jurisdictions. State legislatures are not creating new doctrinal categories for AI. They are extending existing licensing, consumer-protection, and professional-responsibility law to cover AI systems and the entities that deploy them.

# Part III. Global Malpractice Momentum

U.S. courts and state legislatures are not the only forums acting. Comparable enforcement and rule-making is now underway across several other developed legal systems. The convergence is significant for two reasons. First, it reduces the credibility of any defense that AI-driven professional misconduct is a uniquely American problem amenable to a narrow American fix. Second, it forecloses any business strategy that would rely on jurisdictional arbitrage; the regulatory direction is the same on multiple continents simultaneously.

## A. United Kingdom — Ayinde and Al-Haroun (June 2025)

*R (Ayinde) v. London Borough of Haringey* and *Al-Haroun v. Qatar National Bank*, [2025] EWHC 1383 (Admin) (decision of 6 June 2025). In a judgment delivered by Dame Victoria Sharp, President of the King's Bench Division, the High Court of England and Wales considered two cases jointly under the court's Hamid jurisdiction in which AI-generated authorities had been put before the court without verification.

In Al-Haroun, the court's judicial assistants reviewed 45 cited authorities. They found that 18 of the cases did not exist, and that many of the real cases cited contained none of the quotations attributed to them. Justice Sharp wrote that lawyers who fail to verify AI-generated research before submission "risk severe sanction." The court catalogued the available remedies: public admonition, costs orders, contempt proceedings, and referral to the police. The English regulatory architecture (Bar Standards Board Handbook, Solicitors Regulation Authority Code of Conduct) treats the duty to verify as a pre-existing professional obligation, not a new one specific to AI.

The U.K. position aligns closely with the U.S. position emerging in Brigandi and Mata: AI tools are permissible, but the lawyer's duty to verify is non-delegable and pre-existing.

## B. European Union — Regulation (EU) 2024/1689 (the AI Act)

The European Union's Artificial Intelligence Act (Regulation (EU) 2024/1689) entered into force on August 1, 2024. The Act establishes a comprehensive risk-tiered regulatory framework for AI systems within the EU. AI-enabled medical devices, including software as a medical device that incorporates machine learning, are classified as "high-risk AI systems" under Article 6(1) and Annex I. High-risk systems are subject to obligations covering data governance, technical documentation, transparency, human oversight, post-market monitoring, and conformity assessment, in addition to the existing Medical Device Regulation (EU MDR, Reg. 2017/745) requirements.

Most provisions affecting high-risk systems were originally scheduled to enter into force on August 2, 2026. As of April 2026, the European Commission had proposed a "Digital Omnibus" amendment that would extend certain deadlines: standalone high-risk systems to December 2, 2027, and AI embedded in regulated products (including medical devices) to August 2, 2028. The trilogue negotiation on the Annex I architecture failed on April 28, 2026; manufacturers should plan for either timeline.

Maximum penalties under the AI Act reach the higher of €15 million or 3 percent of worldwide annual turnover for general violations, and significantly higher for prohibited-practice violations. Whether enforcement reaches healthcare deployment scenarios at the level seen in U.S. wiretap class actions remains to be seen. The architecture, however, is now in place.

## C. Canada — Ko v. Li (Ontario, 2025)

*Ko v. Li*, 2025 ONSC 2766. In a matrimonial application before the Ontario Superior Court of Justice, applicant's counsel relied on multiple Canadian "court cases" cited in her factum and oral submissions. The presiding judge identified material problems: hyperlinks pointed to unrelated cases; one hyperlink resolved to an HTTP 404 error; another case, when actually examined, reached the opposite conclusion of what counsel had represented. The court ordered counsel to show cause why she should not be cited for contempt of court.

The Ontario decision is doctrinally consistent with the U.S. and U.K. positions. The duty to read cases before submitting them as authority is pre-existing. AI does not change the duty; it changes the failure mode.

## D. Australia and the Cumulative International Picture

Comparable AI-fabricated-citation incidents have been reported in Australian courts, though the formal sanctions architecture there has been less consolidated than in the U.S. and U.K. Charlotin's database, which tracks documented hallucination incidents globally, recorded over 1,300 instances by early 2026 across U.S. federal and state courts, the U.K. High Court, courts in Canada, Australia, the Netherlands, France, and elsewhere. The trend is a one-way ratchet.

For healthcare AI specifically, the international picture is younger but converging. The EU AI Act, AB 3030, AB 489, the Pennsylvania Character Technologies action, the U.K. MHRA's evolving guidance on AI in medical devices, and Health Canada's draft Pre-Market Guidance on Machine Learning-enabled Medical Devices (in consultation as of late 2025) all point toward the same destination: AI in healthcare is regulated as healthcare, with the operator and clinician retaining responsibility, regardless of the AI vendor's marketing about autonomy.

# Part IV. Proleptic Scenario: Case Study 2027-A (Hypothetical)

The following case study is a hypothetical scenario constructed for analytical purposes. It does not describe any actual patient, physician, or pending matter. It is presented to illustrate how the verified doctrinal trajectory in Parts I through III would interact with a specific, but typical, fact pattern under Pennsylvania law.

## Patient profile (hypothetical)

A 49-year-old male presents with a documented history of syringomyelia, dysautonomia, gastroparesis confirmed by gastric-emptying study, and anhidrosis confirmed by quantitative sudomotor axon reflex testing. The patient has a longstanding electronic health record (Epic) note dating from a prior episode of care that characterizes his pain presentation as primarily psychiatric, despite the subsequent objective findings. The treating neurologist orders a Cine-MRI to evaluate progression of the syringomyelia.

## Failure mode (hypothetical Q4 2026)

A cloud-based utilization-management AI deployed by the patient's insurance carrier evaluates the prior-authorization request for the Cine-MRI. The AI weighs the historical psychiatric characterization heavily, hallucinates that prior nerve-block procedures "were definitive" for this patient (a hallucination unsupported by the record), and recommends denial. A medical director rubber-stamps the recommendation in approximately eight seconds, without opening the underlying records, without examining the contradictory objective findings, and in apparent reliance on the AI summary.

The patient's syringomyelia progresses without imaging-guided intervention. The progression results in catastrophic neurological injury that imaging-guided intervention would, on expert testimony, more likely than not have prevented or substantially mitigated.

## Theory of liability

Pennsylvania law provides that medical clinical judgment is non-delegable. A physician who possesses, or has access to, contradictory objective data, and who relies on AI output without examining that data, is liable for the resulting harm under ordinary negligence principles. The eight-second rubber-stamp converts ordinary negligence into a question of willful misconduct or reckless indifference, because the physician possessed (in the EHR available to him) the very contradictory data the AI ignored. The decision was not unconsidered; it was not considered. That distinction matters under Pennsylvania's punitive-damages standard.

California's SB 1120, were the same facts to arise in California, would supply a separate basis for liability by prohibiting medical-necessity decisions based solely on AI output. Pennsylvania has no direct equivalent statute as of May 2026, so the theory in Pennsylvania proceeds under common-law negligence and willful-misconduct doctrine, supported by the standard of care articulated in expert testimony.

# Part V. Damage Modeling: A Defensible Range

This section presents a damage-modeling framework anchored in verified Pennsylvania medical-malpractice verdicts and the Pennsylvania statutory framework governing punitive damages. It does not predict a specific verdict on a specific date. It models a range, with disclosed methodology, into which a hypothetical AI-driven catastrophic-injury case under Pennsylvania law would fall. Plaintiffs may obtain less. Plaintiffs in unusual cases have obtained more. The point of the model is not the point estimate; it is the architecture.

## A. The Pennsylvania Statutory Framework

Pennsylvania does not cap compensatory damages in medical-malpractice cases. The state constitution would not permit such a cap. Punitive damages against an individual physician are capped at 200 percent of compensatory damages (40 P.S. § 1303.505(d)), unless the conduct constitutes intentional misconduct, in which case the cap does not apply. The standard for any award of punitive damages requires conduct that is willful, wanton, or shows reckless indifference to the rights of others. Twenty-five percent of any punitive award is allocated to the Medical Care Availability and Reduction of Error (MCARE) Fund.

Pennsylvania's wrongful-death and survival statutes (42 Pa.C.S. §§ 8301, 8302) supply additional recovery in cases of fatality.

## B. Comparable Verified Pennsylvania Verdicts

| Case | Year | Total Award | Composition |
|---|---|---|---|
| Hagans v. Hosp. of the Univ. of Pa. (Phila. Cty.) | Verdict 2023; upheld 2024 | $183M (largest in PA history) | Birth injury; 45-min delay in C-section; severe neurological injury |
| Heffelfinger v. Shen (Luzerne Cty.) | Feb. 2024 | $11M | Failure to diagnose tongue cancer; $3M comp. + $8M punitive (records destruction) |
| McLaughlin v. McLaughlin (medical malpractice; Phila.) | 2010 | $20.5M | Hospital error; severe brain damage; included $15M punitive (then-record) |
| Bernavage v. Green Ridge Healthcare | Vacated May 2025 | $2.7M punitive vacated | Cautionary case: punitive vacated for procedural reasons ("unfair surprise" |

| Case | Year | Total Award | Composition |
|---|---|---|---|
| | | | amendment) |

These verified anchors illustrate the range. Catastrophic neurological injury cases in Pennsylvania reach mid-eight-figures and, in the most severe cases, nine figures. Failure-to-diagnose cases with documented willful misconduct (in Heffelfinger, the dentist's destruction of records) cross into eight figures. Cases without willful misconduct typically settle below the first-quartile median, often in the low-to-mid seven figures.

## C. Constructing the Range for an AI-Driven Catastrophic Case

A hypothetical case fitting the proleptic scenario in Part IV would, under Pennsylvania law, plausibly produce a verdict in the following defensible range. Each component is anchored to the framework above.

| Component | Conservative | Mid-range | Aggressive |
|---|---|---|---|
| Compensatory: lifetime medical care, lost earnings, pain & suffering | $3M | $8M | $15M |
| Punitive (subject to 200% cap unless intentional misconduct found) | $0 | $8M (1× compensatory) | $30M (2× compensatory cap) |
| Statutory wiretap multiplier (only if Sutter-style facts present) | $0 | $0–$1M | $1M+ (per-recording stacking) |
| **Indicative total exposure** | **≈ $3M** | **≈ $16–17M** | **≈ $46M+** |

These figures are not predictions. They are illustrations of the architecture. The conservative scenario corresponds to a settled case with limited willful-misconduct evidence; the mid-range scenario corresponds to a tried verdict with willful-misconduct findings supporting punitive damages at the 1:1 ratio; the aggressive scenario corresponds to a tried verdict where the conduct supports a finding of intentional misconduct (which lifts the punitive cap entirely, per 40 P.S. § 1303.505(d)) and where statutory wiretap claims are joined.

The honest takeaway is not that any single verdict will land at any specific number. It is that a Pennsylvania case combining catastrophic neurological injury with documented willful misconduct in AI-mediated decision-making sits on a multi-million dollar verdict trajectory, with realistic exposure in the low-to-mid eight figures before extraordinary fact patterns push it higher. That trajectory is supported by published Pennsylvania verdicts and by the existing statutory framework. It does not require a novel theory of AI liability.

# Part VI. Anticipated Plaintiff Theories

The remainder of this paper addresses how plaintiffs' counsel are likely to construct cases of this type. The analysis here is offered in the paper's own voice. It does not include attributed quotations from any actual or simulated attorney. Readers seeking attorney perspective on these theories should consult licensed counsel in the relevant jurisdiction.

## A. The Non-Delegation Theory

Plaintiffs' counsel will argue that clinical judgment is non-delegable, citing the well-established doctrine that the standard of care is what a reasonably prudent clinician would do under the same circumstances. AI output is, under this theory, akin to any other tool: a stethoscope, a laboratory result, an imaging study. The clinician's duty is to exercise judgment about the tool's output, not to ratify it. A physician who signs a denial in eight seconds without opening the chart has not exercised judgment. The non-delegation argument is doctrinally clean and does not require the court to develop new AI-specific law.

## B. The Willful Misconduct Conversion

The pivot from ordinary negligence to willful misconduct (or reckless indifference) is the most consequential element of any plaintiff's theory, because it unlocks punitive damages. Plaintiffs' counsel will argue that a clinician who possessed (or had ready access to) contradictory objective data, and who relied on AI output that was facially in tension with that data, was not negligent in failing to know better; the clinician was reckless in choosing not to look. The architecture of the case is built on the timing: the eight-second decision, documented in system logs, becomes the proof of recklessness. The chart is in the system. The clinician did not open it. The decision is in the system. The AI hallucination is in the system. The expert witness ties the chain together.

## C. The Vendor Co-Defendant Theory

Plaintiffs' counsel will name the AI vendor as co-defendant under product-liability and negligent-design theories. The vendor's marketing claims about accuracy, the absence of guardrails preventing the specific class of hallucination at issue, and the vendor's economic incentive to optimize for token cost rather than clinical accuracy will be discoverable. Vendor internal communications about known failure modes will be discoverable. The vendor's terms of service disclaimers will, in catastrophic-injury cases, be argued unenforceable as against public policy. The Sutter Health complaint already names the AI vendor (Abridge) in the factual narrative; in subsequent cases plaintiffs' counsel will name the vendor as co-defendant from the outset.

## D. The Class Aggregation Theory

For cases with statutory components (wiretap, CIPA, CMIA, license-misrepresentation), plaintiffs' counsel will pursue class certification. The damages model in a certified class is multiplicative, not additive. The Sutter Health complaint is the template: per-recording statutory damages, per-violation statutory damages, and aggregate exposure that scales with deployment volume. A health system that deployed an ambient scribe across 50,000 encounters in six months without consent has potential statutory exposure exceeding $250 million on the federal Wiretap Act alone, before CIPA, CMIA, attorneys' fees, and injunctive relief.

## E. The Pattern-and-Practice Argument

Plaintiffs' counsel will argue that the pattern documented across Brigandi, Sullivan & Cromwell, Character Technologies, and the Charlotin database constitutes industry-wide notice that the conduct at issue is foreseeable and avoidable. The argument is that a defendant who proceeded with cloud AI deployment in a healthcare-clinical-decision context after April 2026 cannot credibly claim the risk was not known. Foreseeability is the gateway to punitive damages, and after a certain volume of public enforcement, the foreseeability question is closed.

# Part VII. Why Cloud AI Fails the Legal Test

The empirical record on enterprise AI performance is now substantial and unflattering. A Carnegie Mellon University study conducted in collaboration with Salesforce and reported in October 2025 found that AI agents fail to reliably complete most office tasks, with failure rates approaching 70 percent in simulated enterprise environments. Gartner's most recent assessment, published in mid-2025 and updated through Q1 2026, projects that more than 40 percent of agentic AI projects will be canceled by the end of 2027 due to escalating costs, unclear ROI, and inadequate risk controls.

RAND Corporation's late-2025 enterprise AI study found that 80.3 percent of enterprise AI projects fail to deliver promised business value, a finding Gartner's April 2026 update directionally confirmed. Carnegie Mellon, RAND, and Gartner are not in agreement on every methodology, but they are in agreement on the direction: enterprise AI is failing at scale.

Worldwide AI spending continues to accelerate notwithstanding. Gartner's January 2026 forecast projected total worldwide AI spending of approximately $2.52 trillion in 2026, up from $1.5 trillion in 2025, with generative AI specifically forecast at $644 billion in 2025. Hyperscaler capital expenditure committed to AI infrastructure for 2026 alone exceeds $450 billion. The gap between deployment volume and demonstrated reliability is the precise condition that creates moral hazard. Juries punish moral hazard.

For healthcare-specific AI, the stakes are sharper. The patient-safety organization ECRI ranked AI chatbot misuse in healthcare as the number-one health technology hazard for 2026. Imperial College London and Microsoft Research have published findings indicating that AI agents paradoxically increase workplace burden as staff are required to supervise and correct the agents' output, undermining the productivity rationale that justified deployment in the first place.

A defendant arguing in a Pennsylvania courtroom in 2027 or 2028 that cloud AI was a reasonable choice for clinical-decision support will face cross-examination armed with all of the foregoing. The cross-examination will not require the plaintiff's expert to know anything special about AI. It will require the expert to read aloud what the industry's own analysts have already published.

# Part VIII. The Enterprise Autonomy Architecture (EAA)

Pocono AI was built for the opposite outcome from the deployment patterns described above. The architecture is not a marketing reframe of cloud AI; it is a different architecture, designed from the outset for environments in which professional liability cannot be transferred to a vendor's terms-of-service page.

## Architectural commitments

- Sentinel Node. A 96GB-VRAM dual-GPU on-premises hardware unit, deployed inside the customer's physical and network perimeter. Inference happens locally. There is no outbound telemetry of patient data during inference, by architectural design rather than by configuration option.
- Mathematical scarcity. The retrieval-augmented generation pipeline (Qdrant, locally hosted) restricts model output to material retrieved from the patient's own sovereign record. The model does not generate from internet-scale training memory in clinical contexts. This is the structural difference that prevents the hallucination-of-novel-content failure mode.
- Immutable audit. Every AI-supported suggestion is logged with the retrieved source chunks, the model version, the prompt, and the human reviewer's sign-off timestamp. The audit log is cryptographically immutable and is the evidentiary substrate that allows a defendant clinician to demonstrate, in litigation, that the AI was a reviewed input, not an unreviewed output.
- Human-in-the-loop as architecture. The system is incapable of generating final output without a human review step. This is not a policy that can be ignored or bypassed under deadline pressure. It is a constraint the software enforces. The Sullivan & Cromwell failure mode (unfollowed policies) cannot occur in this architecture, because the policies are not policies; they are gates.
- Workforce model. Pocono AI's hiring model specifically targets engineers, clinicians, and reviewers who have been sidelined by health, caregiving, or neurodiversity from full-time conventional employment. Roles are structured for capacity and reliable part-time work. The AI handles repetition. Humans retain judgment and accountability. This is documented in detail in the company's published hiring practices and is not a slogan.

## Cost analysis

Capital expenditure for a Sentinel Node is approximately $9,750. Operating cost is local electricity and routine maintenance; there are no per-token cloud charges, no usage-based billing, and no tail-risk exposure to vendor pricing changes. By comparison, the conservative damages model in Part V begins at approximately $3 million in indicative total exposure for a single catastrophic-injury verdict against a cloud-AI-using defendant. The math, even before counting class-action multipliers and statutory damages stacking, is not close.

An organization that finds the Sentinel Node CapEx unaffordable should reconsider whether it can afford the deployment of cloud AI at all in clinical-decision contexts. The architecture exists. The choice is no longer between cloud AI and no AI; the choice is between cloud AI and architecturally sovereign AI. A defendant in a 2028 trial who selected the cloud option will be asked to explain why.

## Part IX. Conclusion: Custody Is Liability

Four cases, three jurisdictions, thirty days. Over 1,300 documented hallucination incidents globally. A Carnegie Mellon failure rate near 70 percent. A Gartner project-cancellation forecast above 40 percent. A Pennsylvania verdict architecture that produces eight-figure exposure on standard fact patterns and nine-figure exposure on extraordinary ones. An EU regulatory framework converging on the same direction with maximum penalties of 3 percent of global turnover. A U.K. High Court warning of "severe sanction" up to and including referral to the police. A Pennsylvania state Board of Medicine prosecuting an AI chatbot under the same statute that has disciplined unlicensed humans for over a century.

The era of renting clinical intelligence from external vendors and signing one's professional license to the output is ending. It is not ending because regulators have been slow; regulators have been moving for two years. It is ending because the foreseeability question is closing in real time, and once it closes, every subsequent rubber-stamp becomes evidence of recklessness.

The defensible architecture is physical custody of model weights, locally retrieval-grounded inference, immutable audit, and provable human-in-the-loop review. The Sentinel Node operationalizes that architecture in a unit a solo physician can purchase, install, and verify without specialized engineering staff. The cost of doing this is approximately five figures. The cost of not doing this, when the first significant verdict lands, will be substantially greater.

When the first AI-driven medical-malpractice verdict in Pennsylvania crosses eight figures, the marketplace will not respond by adopting better cloud prompts. It will respond by demanding what was always available: AI that does not phone home with the patient's record, AI that cannot generate beyond what the record actually contains, AI that has signed a real audit trail before any human signs a real chart note. That is the product. That is the defense. That is the company.

# Appendix A. Verification Log

Each citation in this paper has been verified against primary or authoritative secondary sources as of May 10, 2026. Where a v2 version of this paper contained errors, those errors are documented and corrected here. Court filings and regulatory documents continue to evolve; readers contemplating reliance on any specific citation should verify it against the current public docket.

| Claim | Primary or Authoritative Source | Verification Date |
|---|---|---|
| Brigandi/Murphy $110K Oregon sanction; Op. filed Apr. 4, 2026 (D. Or.); "notorious outlier" quote; vineyard dismissed with prejudice | ABA Journal (Apr. 17, 2026 article); Rogue Valley Times; KGW.com; Yahoo News (April 2026); Ethics Reporter case profile | May 10, 2026 |
| Sullivan & Cromwell letter to Chief Judge Martin Glenn dated Apr. 18, 2026; original filing Apr. 9, 2026; errors caught by Boies Schiller Flexner | Bloomberg Law; Bloomberg News; Reuters (via PYMNTS coverage); Original Jurisdiction (David Lat); Legal Cheek | May 10, 2026 |
| Commonwealth v. Character Technologies, No. 220 MD 2026 (Pa. Commw. Ct.), filed May 1, 2026; Medical Practice Act, 63 P.S. §§ 422.1 et seq. | Pa. Commw. Ct. docket; ppc.land coverage; Reuters (May 5, 2026); CBS News; The Next Web; Regulatory Oversight (Troutman Pepper Locke analysis) | May 10, 2026 |
| Washington v. Sutter Health, et al., No. 4:26-cv-03012-KAW (N.D. Cal., filed Apr. 8, 2026); Sutter and MemorialCare named; CIPA, CMIA, Wiretap Act claims | N.D. Cal. docket; HIPAA Journal; TechTarget; Yahoo News; complaint excerpt (ismg-cdn) | May 10, 2026 |
| California AB 3030, eff. Jan. 1, 2025; Cal. Health & Safety Code § 1339.75 | California Legislative Counsel (leginfo.legislature.ca.gov); Medical Board of California guidance; multiple law-firm advisories | May 10, 2026 |
| California AB 489, signed Oct. 11, 2025; effective Jan. 1, 2026; Cal. Bus. & Prof. Code § 4999.8 et seq. | California Legislative Counsel (leginfo.legislature.ca.gov); Hintze Law analysis; Akerman LLP analysis | May 10, 2026 |

| Claim | Primary or Authoritative Source | Verification Date |
|---|---|---|
| Carnegie Mellon / Salesforce study, ~70% AI agent failure rate on office tasks (Oct. 2025) — NOT Gartner, as v2 incorrectly attributed | Carnegie Mellon University news / WorkersCompensation.com coverage; cross-referenced via Mighty Blog and multiple secondary analyses | May 10, 2026 |
| Gartner: >40% agentic AI projects canceled by end of 2027; total AI spend $2.52T in 2026 (44% YoY) | Gartner press release (Jan. 15, 2026); Gartner press release (Aug. 26, 2025); secondary analyses (TheStreet; Yahoo Finance) | May 10, 2026 |
| RAND ~80.3% enterprise AI project failure rate (late 2025); Gartner April 2026 confirmation | MyBusinessFuture analysis citing RAND 2025 and Gartner Apr. 7, 2026 update | May 10, 2026 |
| Hagans v. HUP — $183M PA medical-malpractice verdict; upheld Feb. 2024 | Daily Pennsylvanian (Feb. 6, 2024); Expert Institute coverage | May 10, 2026 |
| Heffelfinger v. Shen, No. CV-2020-008443 (C.P. Luzerne Feb. 23, 2024) — $11M verdict including $8M punitive | Lawsuit Information Center; Duane Morris alert | May 10, 2026 |
| PA punitive damages cap: 200% of compensatory; 40 P.S. § 1303.505(d) (MCARE Act) | Pennsylvania Statutes; Youman & Caputo analysis; Lupetin & Unatin; Mattiacci Law | May 10, 2026 |
| R (Ayinde) v. Haringey & Al-Haroun v. Qatar National Bank, [2025] EWHC 1383 (Admin) (June 6, 2025); Dame Victoria Sharp | TechCrunch (June 7, 2025); Appleby Global; Counsel Magazine; International Bar Association | May 10, 2026 |
| Ko v. Li, 2025 ONSC 2766 (Ont. S.C.J.) — Canadian contempt sanction for AI-fabricated cases | BCLP (Bryan Cave Leighton Paisner) analysis; Ontario Superior Court records | May 10, 2026 |
| EU AI Act, Reg. (EU) 2024/1689; in force Aug. 1, 2024; high-risk systems originally Aug. 2, 2026 (Digital Omnibus extensions proposed) | Official Journal of the European Union; White & Case alert; ComplyLoft analysis; MedDeviceGuide | May 10, 2026 |

| Claim | Primary or Authoritative Source | Verification Date |
|---|---|---|
| Damien Charlotin AI Hallucination Cases Database (>1,300 documented incidents as of early 2026) | Charlotin database (damiencharlotin.com); coverage in LA Times, Volokh Conspiracy, 404 Media | May 10, 2026 |
| ECRI #1 health technology hazard 2026: AI chatbot misuse in healthcare | ECRI Institute; secondary coverage (The Next Week / TNW) | May 10, 2026 |
| Mata v. Avianca, 678 F. Supp. 3d 443 (S.D.N.Y. 2023); $5,000-per-attorney sanction | FindLaw; Westlaw; Justia; ACC analysis; Seyfarth Shaw analysis | May 10, 2026 |

## Errata from v2

v2 of this paper, dated May 10, 2026 (initial release), contained the following errors corrected in v3:

(1) Brigandi/Murphy opinion date stated as April 17, 2026; corrected to April 4, 2026 (the April 17 date was the ABA Journal article, not the opinion).

(2) Commonwealth v. Character Technologies filing date stated as May 5, 2026; corrected to May 1, 2026 (May 5 was the public announcement). Petitioner correctly identified as Pennsylvania Department of State, State Board of Medicine. Docket number 220 MD 2026 added.

(3) Washington v. Sutter Health corrected to include MemorialCare as co-defendant. Three named plaintiffs (Washington, Gueretta, Matulic) added.

(4) California AB 489 stated as enacted October 11, 2025 with ambiguity about effective date; clarified as signed October 11, 2025, effective January 1, 2026.

(5) 70% AI failure rate misattributed to Gartner; corrected to Carnegie Mellon University / Salesforce study (October 2025). Gartner's actual relevant prediction (>40% agentic AI projects canceled by end of 2027) included separately.

(6) AI spending figures of $86.4B (2025) → $206.5B (2026) removed as unverified; replaced with verified Gartner figures: total AI $1.5T (2025) → $2.52T (2026); generative AI specifically $644B (2025).

(7) Specific $25M+ verdict prediction by late 2027 in Pennsylvania removed. Replaced with damage-modeling framework anchored in verified comparable Pennsylvania verdicts and

the 200% punitive cap under 40 P.S. § 1303.505(d), presenting a defensible range rather than a point estimate.

# Appendix B. Methodology and Limitations

This paper relies on three categories of source. First, primary judicial and regulatory sources: court dockets, slip opinions, statutes, and regulatory filings, accessed through public docket systems and official legislative websites. Second, authoritative secondary sources: legal-industry publications (ABA Journal, Bloomberg Law, Reuters), law-firm advisories (Duane Morris, White & Case, Akerman, Hintze, Bryan Cave Leighton Paisner, and others), and trade press with recognized accuracy in the relevant jurisdictions. Third, where empirical data is cited (Carnegie Mellon, RAND, Gartner, ECRI), the original organizational publication or its closest faithful summary is the source.

Limitations. This paper is not legal advice. Court dockets continue to evolve; an opinion cited as final today may be appealed, vacated, or modified tomorrow. Statutes are amended. Regulatory deadlines are extended (the EU AI Act timeline is itself in flux as of this writing). Damage-modeling figures are illustrative and depend on facts not present in any actual case. The paper's hypothetical scenario in Part IV is constructed for analytical purposes and does not describe any real patient, clinician, or pending matter.

Most importantly, the paper does not predict any specific verdict on any specific date. The damage-modeling framework presents a defensible range based on verified comparable verdicts and the existing Pennsylvania statutory architecture. Plaintiffs may obtain less. Plaintiffs in unusual cases have obtained more. The verdict trajectory described herein is supported by the empirical record; the timing of the first significant verdict is unknowable in advance.

Readers contemplating litigation, compliance restructuring, or commercial action in reliance on any portion of this paper should consult licensed counsel in the relevant jurisdiction. Pocono AI, LLC and the author do not undertake to update this paper as developments occur, but anticipate doing so on a periodic basis.

*— End of Document —*